

CS 7650: Natural Language Processing (Fall 2025)

Problem Set 1

Instructor: Wei Xu

TAs: Rohan Phadnis, Joseph Thomas, Jerry Lou Zheng, Duong Le

Course website: https://cocoxu.github.io/CS7650_fall2025/

Gradescope: <https://www.gradescope.com/courses/1086056>

Due: Friday, September 5, 11:59 PM ET

1 Logistic vs Softmax

1. (2 pts) Recall the Logistic and Softmax functions

$$P_{Logistic}(y = 1|\mathbf{x}) = \frac{e^{\mathbf{w}^T \mathbf{x}}}{1 + e^{\mathbf{w}^T \mathbf{x}}}$$

$$P_{Softmax}(y|\mathbf{x}) = \frac{e^{\mathbf{w}_y^T \mathbf{x}}}{\sum_{y' \in \mathcal{Y}} e^{\mathbf{w}_{y'}^T \mathbf{x}}}$$

Given $\mathcal{Y} = \{0, 1\}$, what should be the value of \mathbf{w} such that $P_{Logistic}(y|\mathbf{x}) = P_{Softmax}(y|\mathbf{x}) \ \forall \ y \in \mathcal{Y}$? Show your work.

Hint: Expand the summation term and think about \mathbf{w} in terms of \mathbf{w}_0 and \mathbf{w}_1 . \mathbf{w}_0 and \mathbf{w}_1 are the weight vectors corresponding to the negative ($y = 0$) and positive ($y = 1$) classes, respectively.

2. (2 pts) Recall that the Softmax function is a generalization of the logistic sigmoid for multiclass classification. In practice, machine learning software such as PyTorch uses a Softmax implementation for both binary and multiclass classification. Also recall that the Softmax function produces a vector output $\mathbf{z} \in \mathbb{R}^{|\mathcal{Y}|}$ and the logistic function a single scalar value z , representing class probabilities. Write the equation for a decision rule to produce \hat{y} from the Softmax function in the binary case (when $\mathcal{Y} = \{0, 1\}$; you can break ties arbitrarily). Write the decision rule to produce \hat{y} from the logistic function. Compare the two rules. How are they similar and/or different? (1-2 sentences).

2 Multiclass Naive Bayes with Bag of Words

A movie studio wants to check if their movie is being received well on social media. In particular, they want to use a Naive Bayes algorithm to automatically classify if a certain review is Positive, Neutral, or Negative. These reviews have already been passed through a feature function which returned the following key features based on the count of certain words used to describe the movie. A set of these reviews have been sampled and each were classified based on their feature vectors. The data collected so far is given in the table below:

S.No.	interesting	terrific	awful	strange	sad	confusing	amazing	amusing	Y
1	1	0	0	1	1	0	0	0	Positive
2	1	0	1	0	0	1	1	1	Neutral
3	0	0	1	1	1	1	0	1	Negative
4	1	0	1	0	0	1	0	0	Neutral
5	0	1	0	1	1	0	0	0	Positive
6	0	1	0	0	0	0	1	0	Neutral
7	0	0	1	1	0	0	1	1	Positive
8	1	0	1	1	1	1	0	0	Negative

- a. (1 pt) What is the probability θ_y of each label $y \in \{\text{Positive, Neutral, Negative}\}$?
- b. (3 pts) The parameter $\phi_{y,j}$ is the probability of a token j appearing with label y . It is defined by the following equation, where V is the size of the vocabulary set:

$$\phi_{y,j} = \frac{\text{count}(y, j)}{\sum_{j'=1}^V \text{count}(y, j')}$$

The probability of a count of words x and a label y is defined as follows. Here, $\text{count}(y, j)$ represents the frequency of word j appearing with label y over all data points.

$$p(x, y; \theta, \phi) = p(y; \theta) \cdot p(x|y; \phi) = p(y; \theta) \prod_{j=1}^V \phi_{y,j}^{x_j}$$

Find the most likely label \hat{y} for the following word counts vector $x = (1, 0, 0, 1, 1, 0, 1, 0)$ using $\hat{y} = \text{argmax}_y \log p(x, y; \theta; \phi)$. Show final log (base-10) probabilities for each label rounded to 3 decimals. Treat $\log(0)$ as $-\infty$. (Hint: read the above provided equations carefully; read more about *binary multinomial naive Bayes* in Jurafsky & Martin Chapter 4, as well as Hiroshi Shimodaira's note - https://cocoxu.github.io/CS7650_fall2025/slides/Shimodaira_note07.pdf.)

- c. (3 pts) When calculating argmax_y , if $\phi_{y,j} = 0$ for a label-word pair, the label y is no longer considered. This is an issue, especially for smaller datasets where a feature may not be present in all documents for a certain label. One approach to mitigating this high variance is to smooth the probabilities. Using add-1 smoothing, which redefines $\phi_{y,j}$, again find the most likely label \hat{y} for the following word counts vector $x = (1, 0, 0, 1, 1, 0, 1, 0)$ using $\hat{y} = \text{argmax}_y \log p(x, y; \theta; \phi)$. Make sure to show final log probabilities.

$$\text{add-1 smoothing: } \phi_{y,j} = \frac{1 + \text{count}(y, j)}{V + \sum_{j'=1}^V \text{count}(y, j')}$$

3 Perceptron: Linear Separability and Weight Scaling

- (a) (**2 pts**) Suppose we have the following data representing the XOR function:

Evidently, the data is not linearly separable. Therefore the perceptron algorithm will not be able to learn a classifier for XOR, based on this data.

x_1	x_2	$f(x_1, x_2)$
0	0	-1
0	1	1
1	0	1
1	1	-1

Table 1: XOR function data

However, we can add a 3^{rd} dimension/feature to each input such that the data becomes linearly separable. If we add $(1, 0, 0, 1)$ to the 3^{rd} dimension (x_3) of the four data points in order, will the new data be linearly separable? Assume 0 is the threshold for classification. Justify your answer.

Does the ability to add a 3^{rd} dimension indicate that the perceptron algorithm is capable of learning the XOR function? Why or why not?

- (b) (**2 pts**) Suppose we have a trained Perceptron with parameters (W, b) . If we scale W by a positive constant factor c , will the new set of weights produce the exact same prediction for all the test data? Assume the threshold for classification is 0. Justify your answer.
- (c) (**2 pts**) With the same setting as (b), this time we translate W by a positive constant factor c (add c to each element of W), will the new set of weights produce the exact same prediction for all the test data? Justify your answer.

4 Feedforward Neural Network

[Eisenstein Chapter 3 Problem 4] (2 pts) In Question 3, we tried to design a perceptron architecture in order to learn the XOR function represented by Table 1. Now, we want you to design a feedforward neural network to compute the XOR function. This can be done in several different ways, so make sure you provide ample description of your design!

Use a single output node and **specify** the activation function you choose for it. Also use a single hidden layer with ReLU activation function. Describe all weights and offsets (bias terms) and ensure you design your network in accordance with Table 1. You may draw a diagram similar to the one shown below (or just indicate the weights).

(Hint: In class, we discussed a neural network design that solves the XOR problem using ***tanh*** activation functions.)

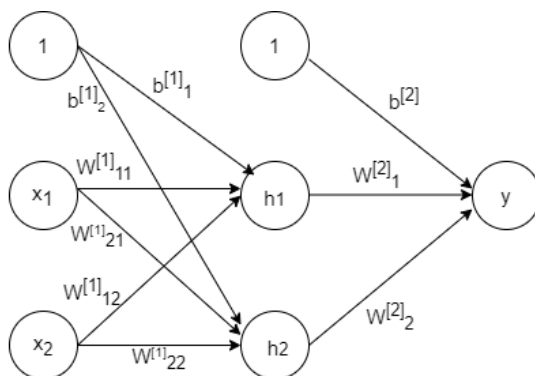


Figure 1: FFNN for XOR

5 Dead Neurons

[Eisenstein Chapter 3 Problem 8] The ReLU activation function can lead to “dead neurons”, which can never be activated on any input. Consider a feedforward neural network with a single hidden layer and ReLU nonlinearity, assuming a binary input vector, $\mathbf{x} \in \{0, 1\}^D$ and scalar output y :

$$\begin{aligned} z_i &= \text{ReLU}(\theta_i^{(x \rightarrow z)} \cdot \mathbf{x} + b_i) \\ \mathbf{y} &= \theta^{(z \rightarrow y)} \cdot \mathbf{z} \end{aligned}$$

Assume the above function is optimized to minimize a loss function (e.g., mean squared error) using stochastic gradient descent.

1. **(2 pts)** Under what condition is node z_i “dead”? Your answer should be expressed in terms of the parameters $\theta_i^{(x \rightarrow z)}$ and b_i .
2. **(2 pts)** Suppose that the gradient of the loss on a given instance is $\frac{\partial l}{\partial y} = 1$. Derive gradients $\frac{\partial l}{\partial b_i}$ and $\frac{\partial l}{\partial \theta_{j,i}^{(x \rightarrow z)}}$ for such an instance.
3. **(2 pts)** Using your answers to the previous two parts, explain why a “dead” neuron can never be brought back to life during gradient-based learning.